DeepSeek-R1-Distill-Llama-70B
TP=4, PP=3, Input=1024, Output=512
Instance=3x g5.12xlarge, GPU=A10G
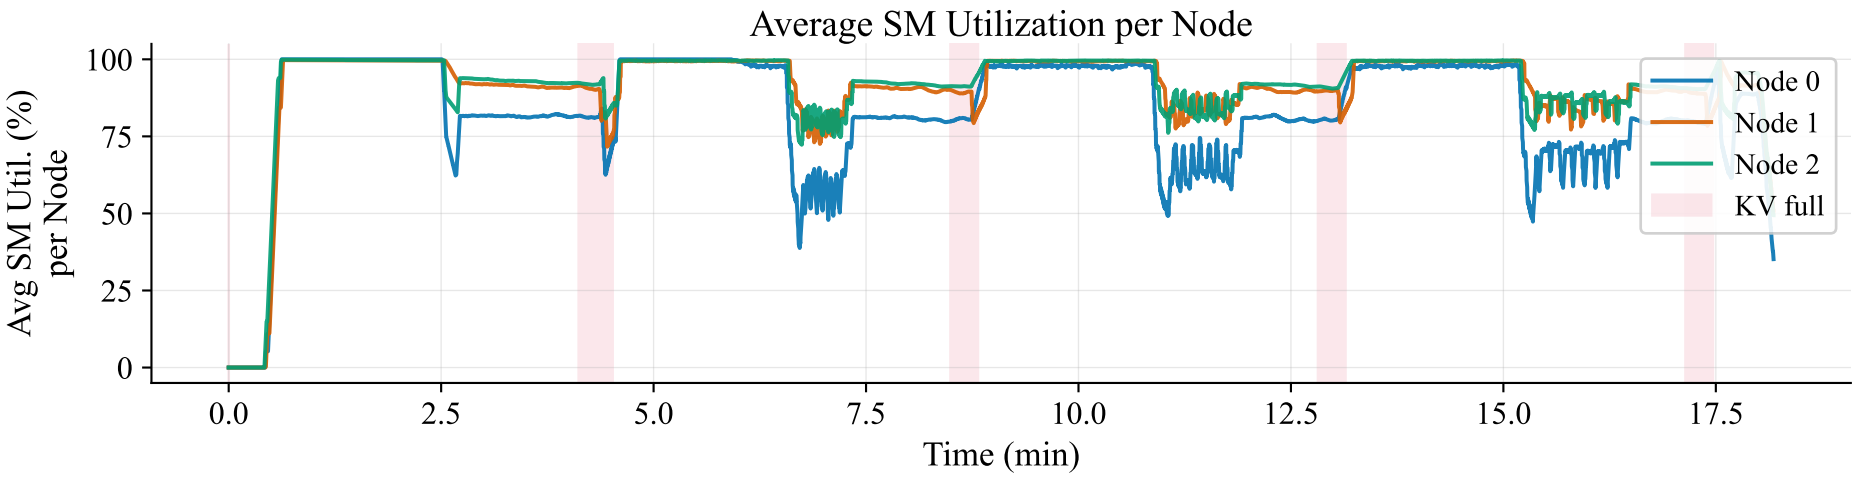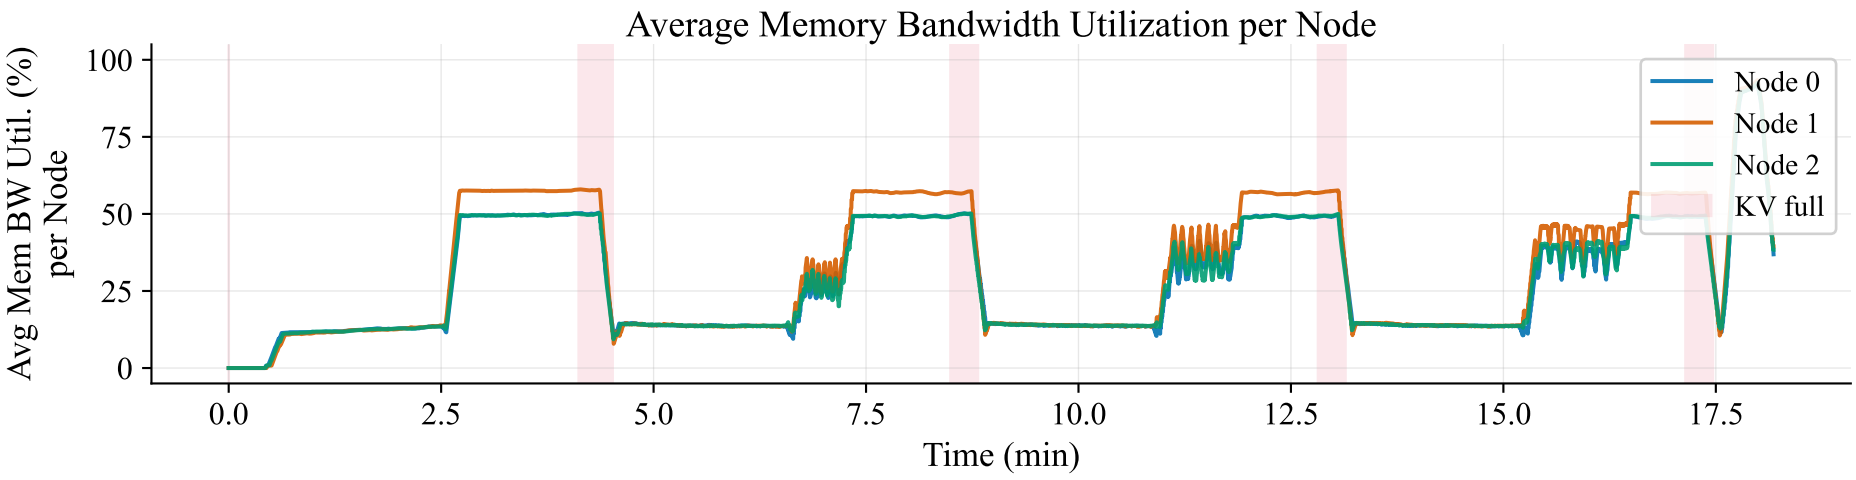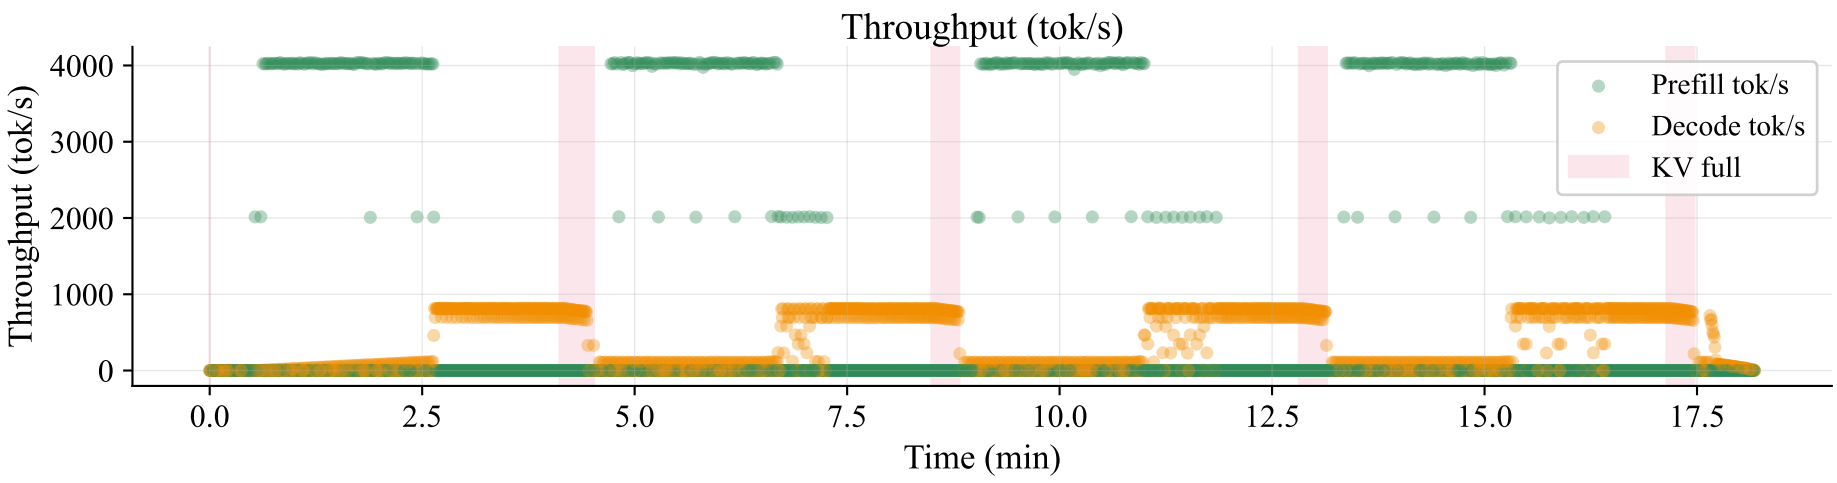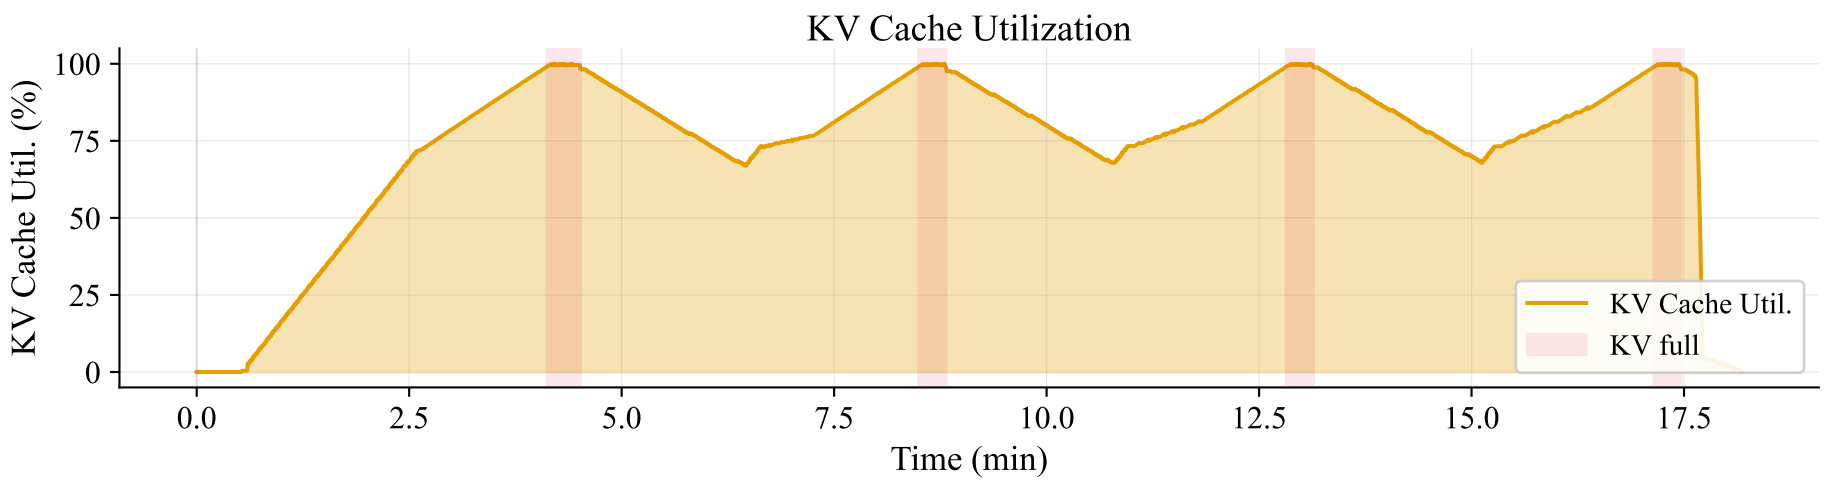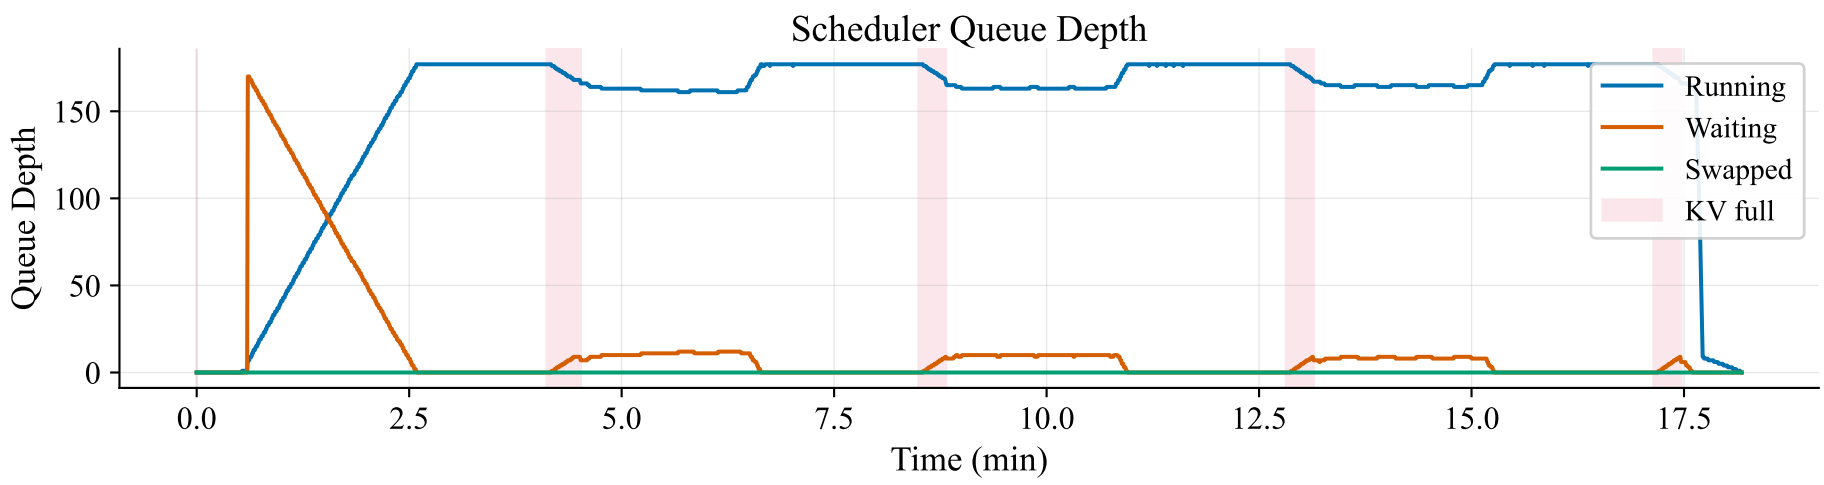
Throughput		Cost & Time		Benchmark Config		GPU Utilization (avg)	
Prefill tok/s	683.9	Elapsed	17m 40s	Mode	rate-limited	SM util	88%
Decode tok/s	341.9	Cost	$3.62	Requests	708	Mem BW util	30%
Requests/sec	0.668	Tok/dollar	300,531	Preemptions	37	Mem util	88%
		$/1M tok	3.33			Nodes	3