DeepSeek-R1-Distill-Llama-70B
TP=4, PP=3, Input=1024, Output=512
Instance=3x g5.12xlarge, GPU=A10G
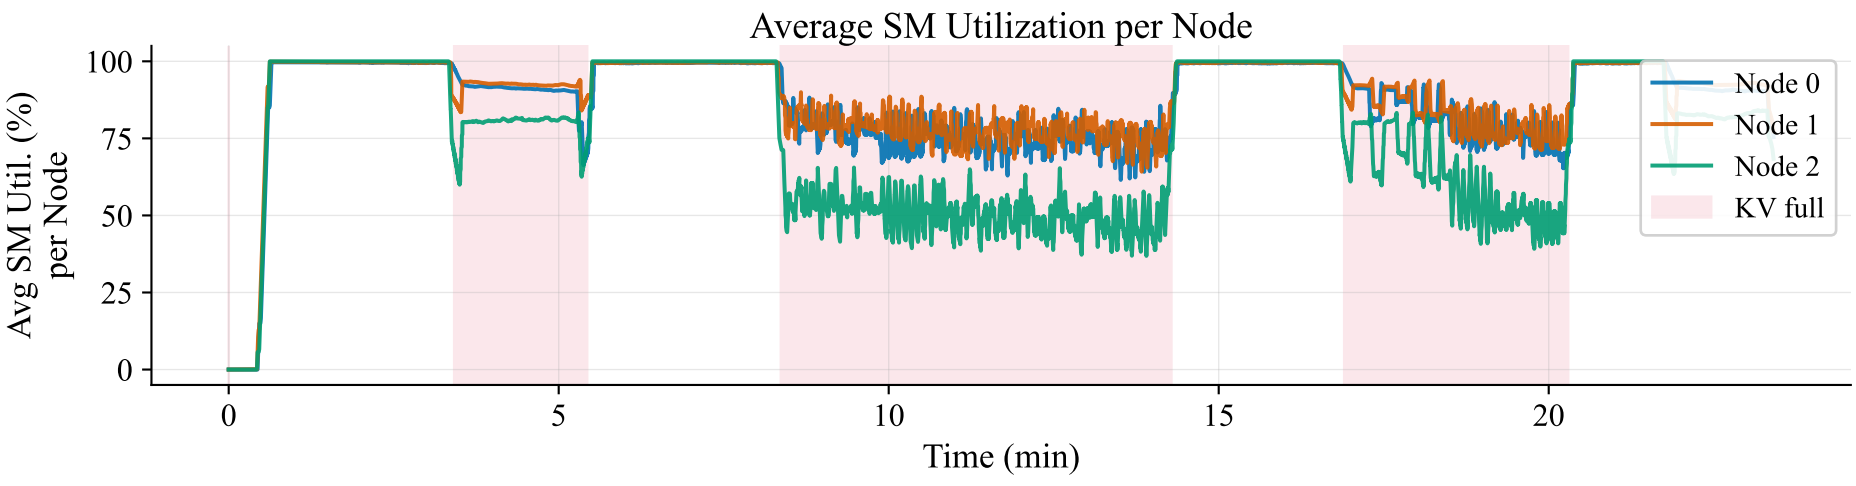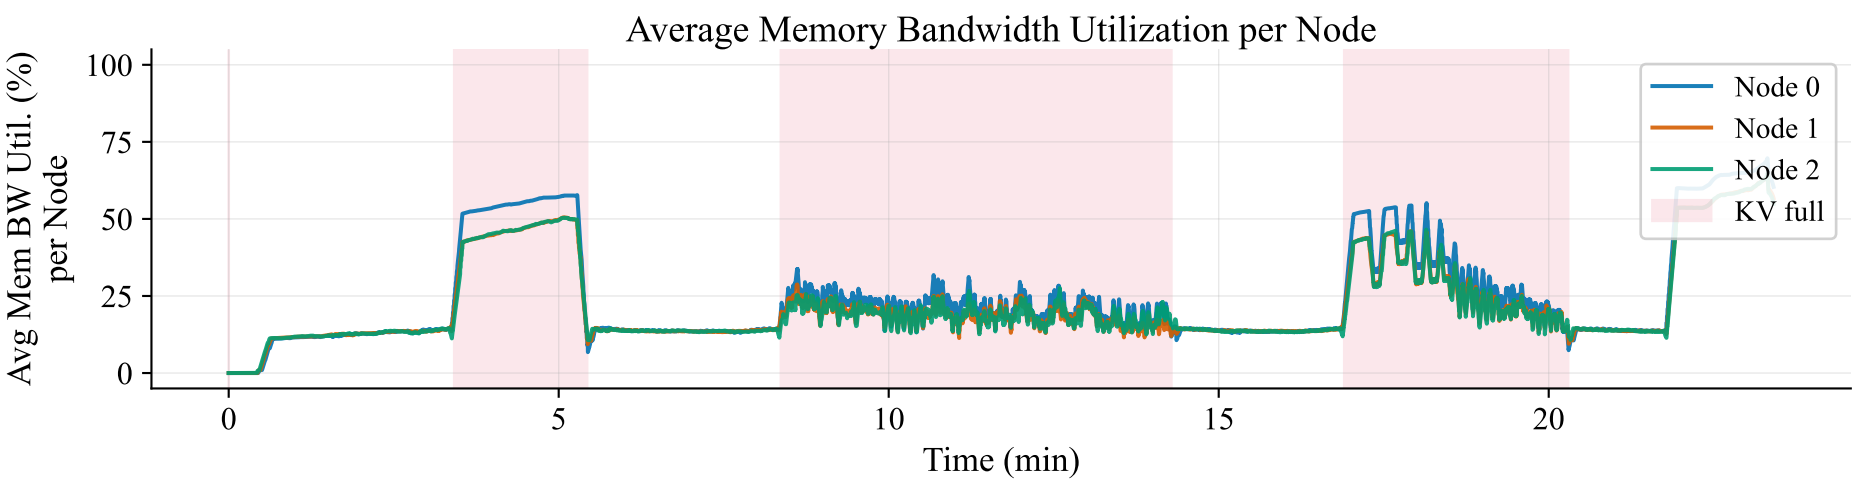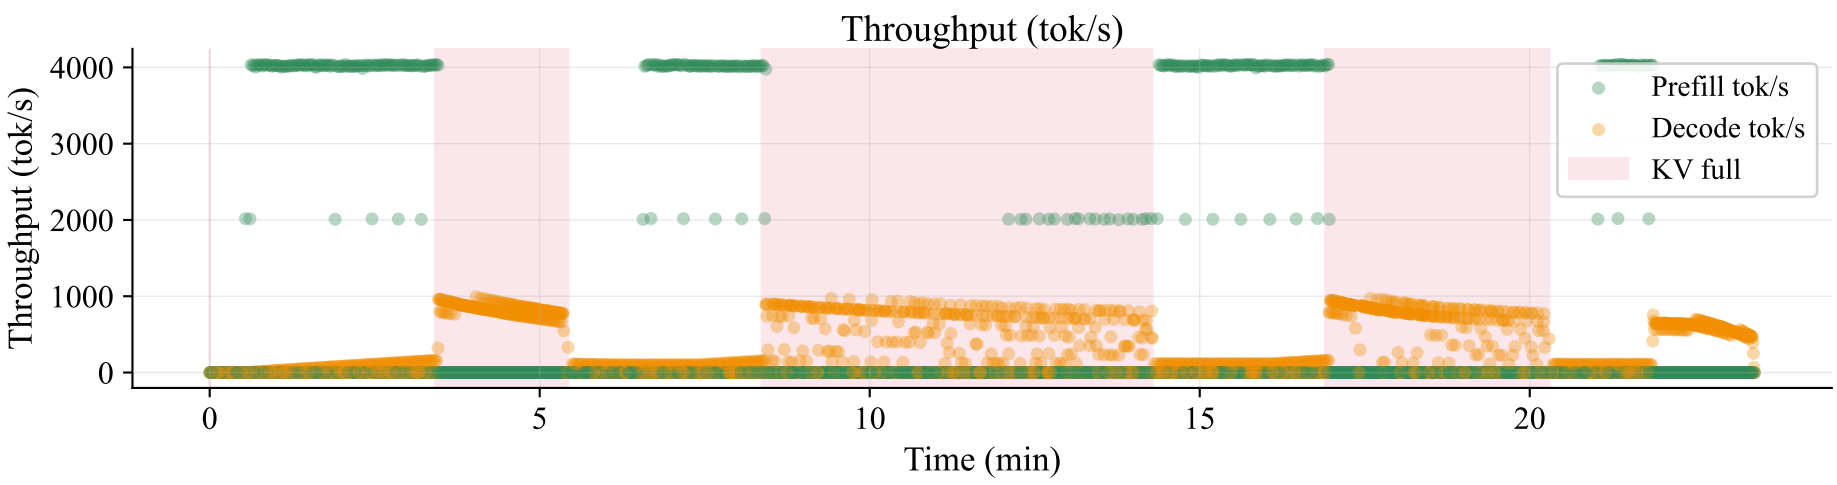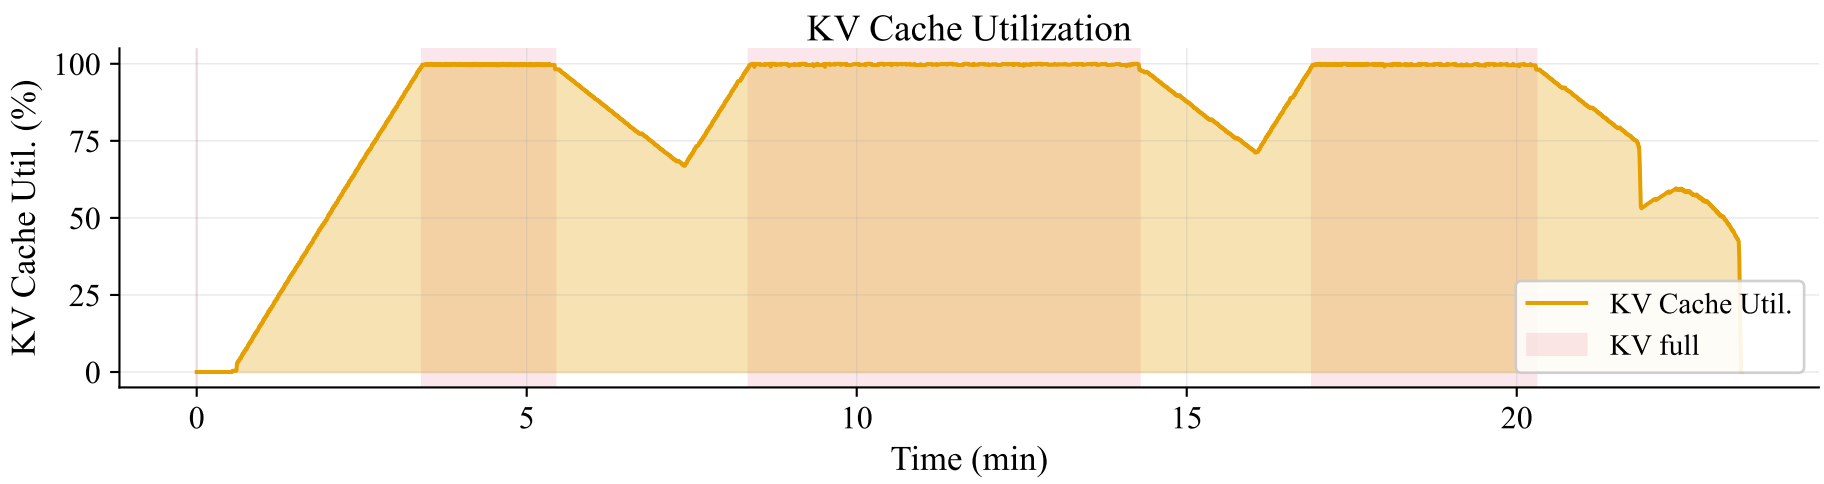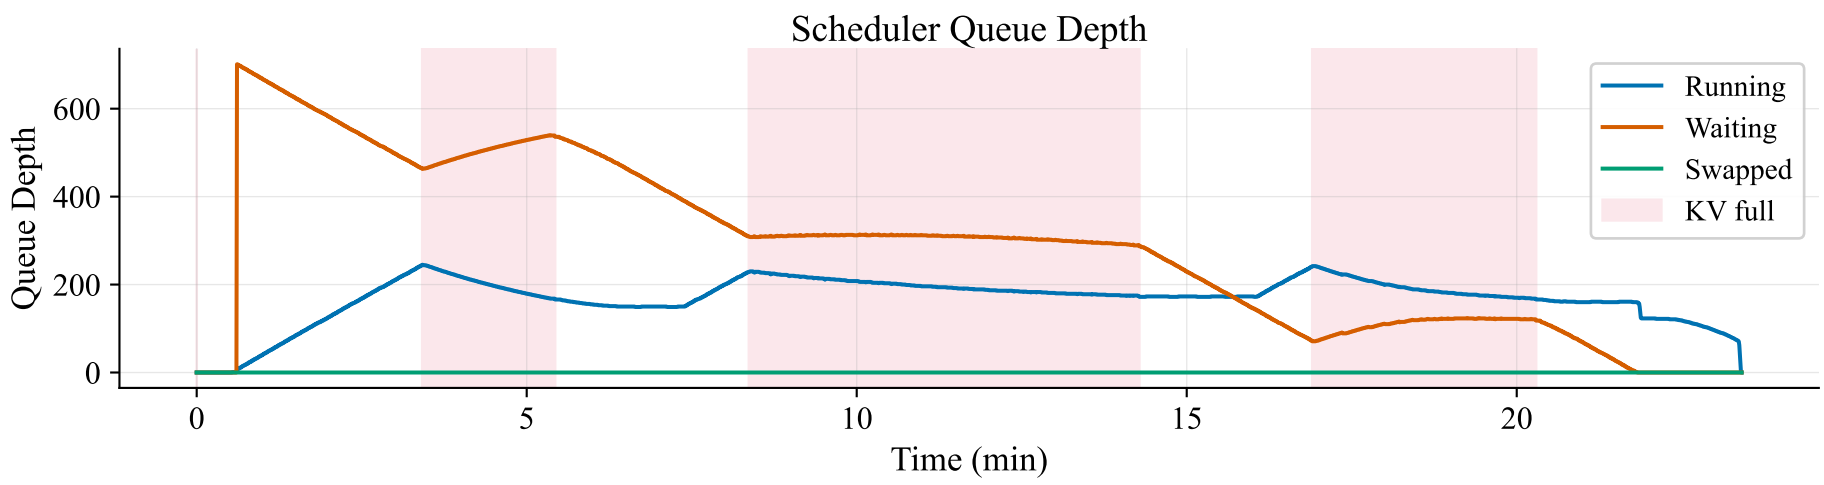
Throughput		Cost & Time		Benchmark Config		GPU Utilization (avg)	
Prefill tok/s	527.7	Elapsed	22m 54s	Mode	send-all	SM util	84%
Decode tok/s	263.8	Cost	$4.69	Requests	708	Mem BW util	23%
Requests/sec	0.515	Tok/dollar	231,887	Preemptions	268	Mem util	88%
		$/1M tok	4.31			Nodes	3