DeepSeek-R1-Distill-Llama-70B
TP=4, PP=3, Input=1024, Output=512
Instance=3x g5.12xlarge, GPU=A10G
Elapsed=1060.1s, GPU Time=1090.8s

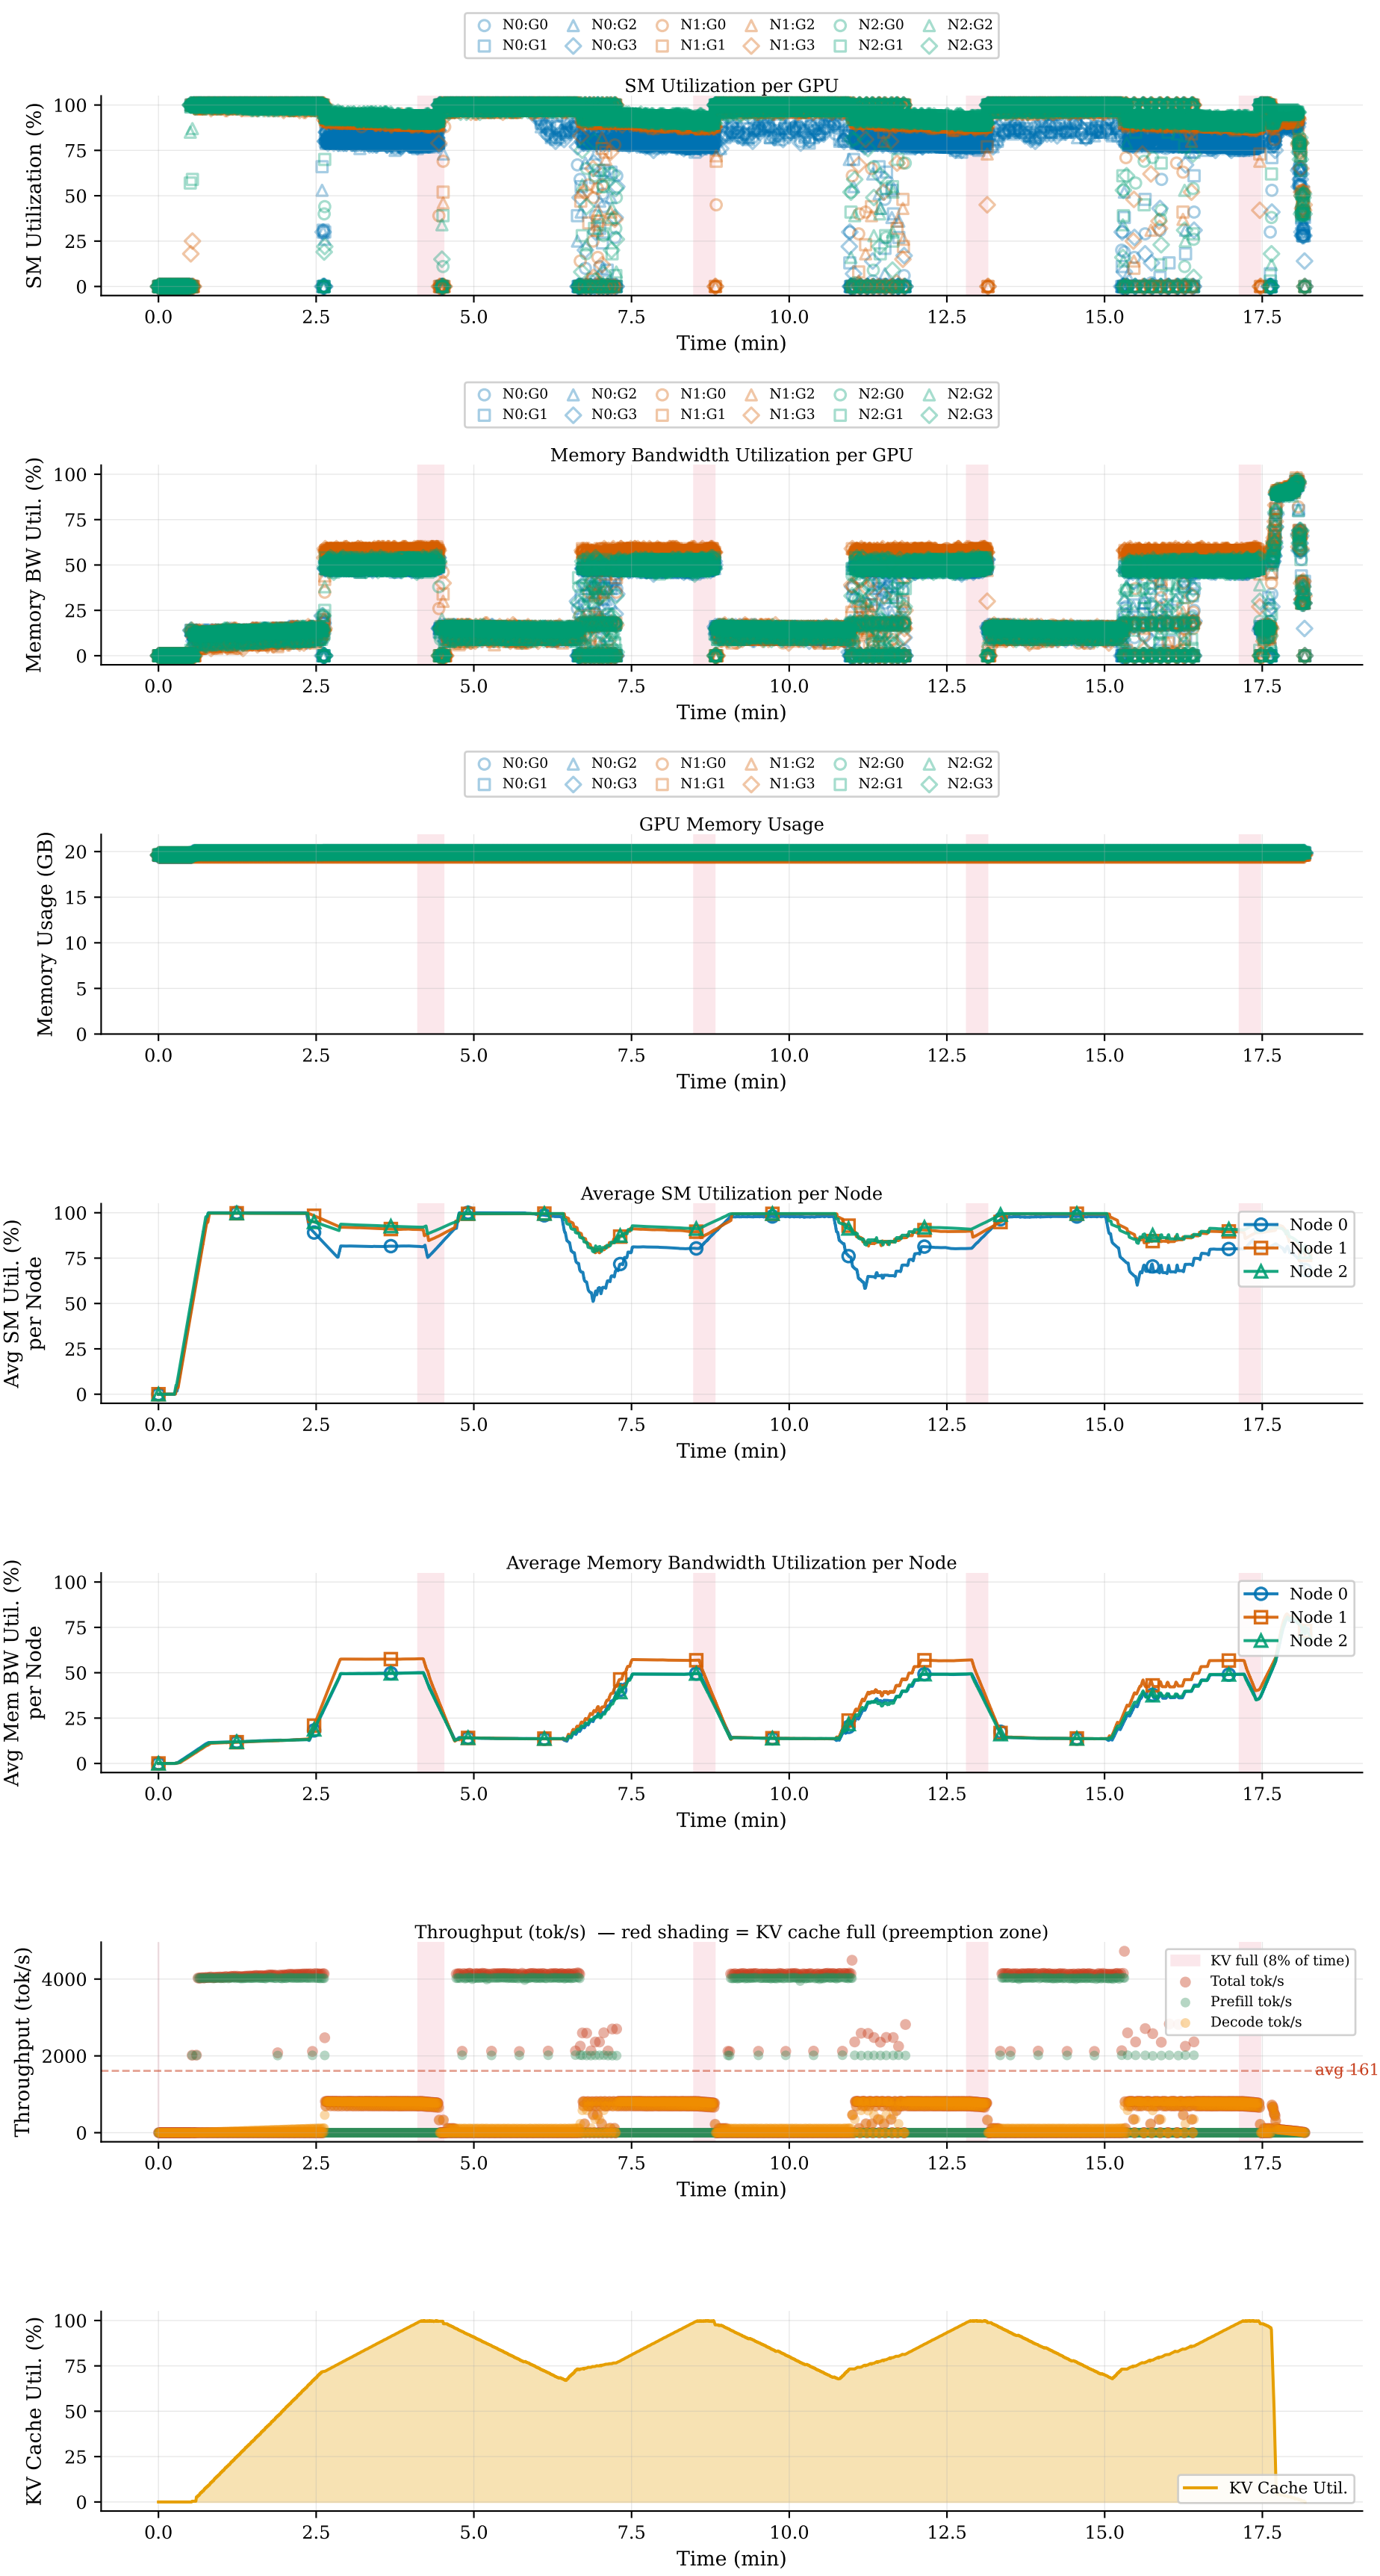

Experiment Summary

Throughput		Cost & Time		Benchmark Config		GPU Utilization (avg)	
Total tok/s:	**1,025.8**	Elapsed:	**17m 40s**	Mode:	**rate-limited**	SM util:	**88.3%**
Prefill tok/s:	**683.9**	Cost:	**$3.6186**	Concurrency:	**177**	Mem BW util:	**29.8%**
Decode tok/s:	**341.9**	Tok/dollar:	**300,531**	Requests:	**708**	Mem util:	**88.0%**
Requests/sec:	**0.668**	$/1M tok:	**3.3275**	Preemptions:	**37**	Nodes:	**3**