DeepSeek-R1-Distill-Llama-70B
TP=4, PP=3, Input=1024, Output=512
Instance=3x g5.12xlarge, GPU=A10G
Elapsed=1373.9s, GPU Time=1404.4s

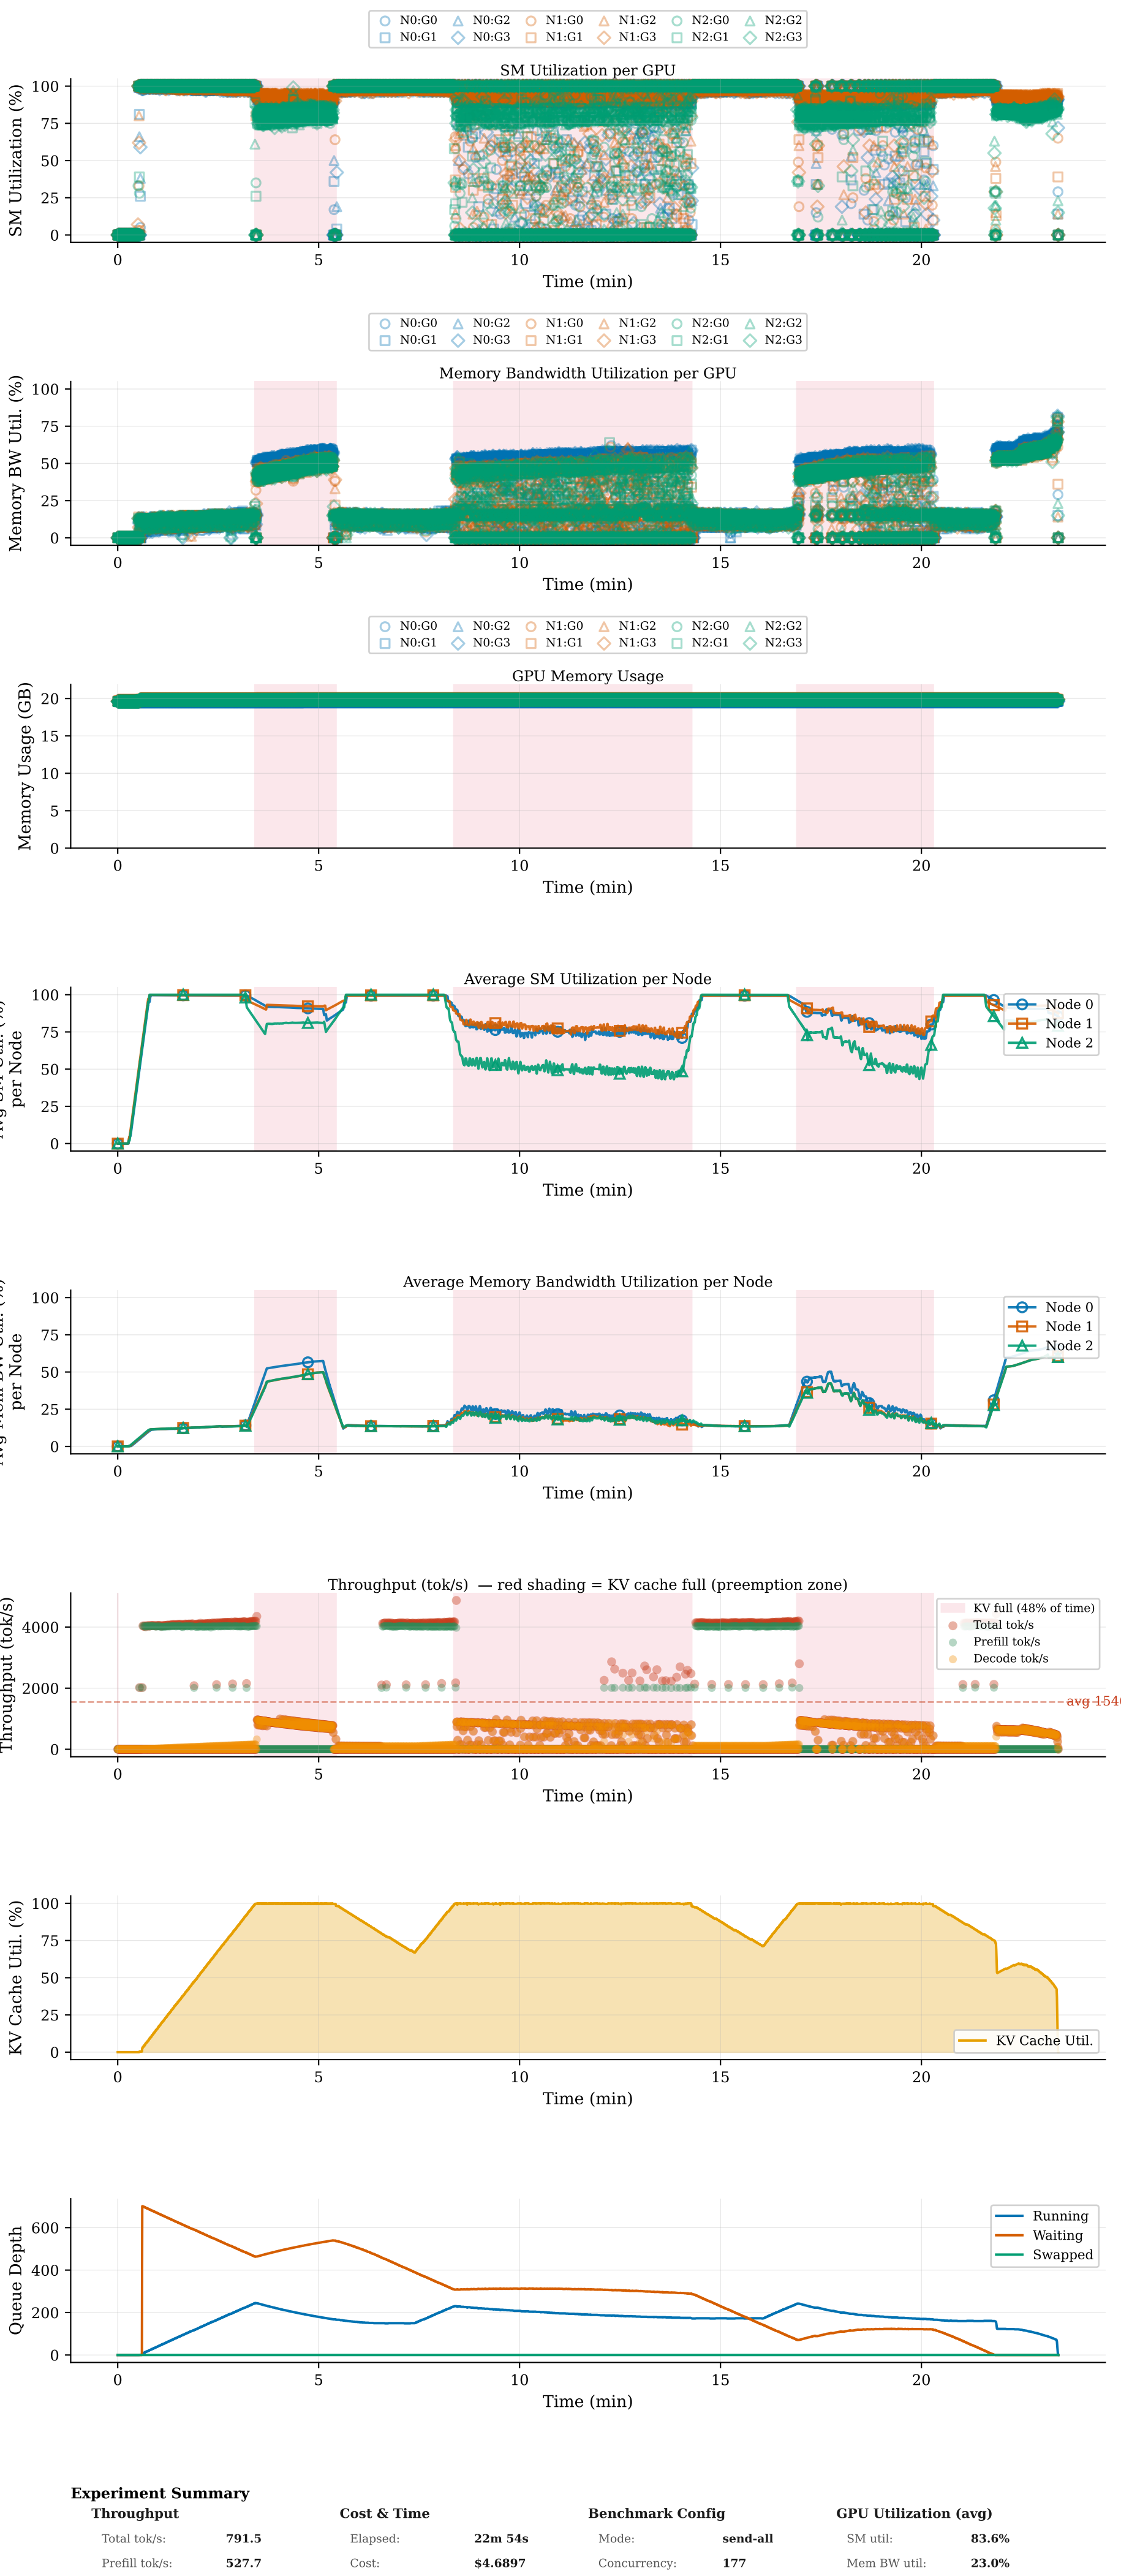

Experiment Summary

Throughput		Cost & Time		Benchmark Config		GPU Utilization (avg)	
Total tok/s:	**791.5**	Elapsed:	**22m 54s**	Mode:	**send-all**	SM util:	**83.6%**
Prefill tok/s:	**527.7**	Cost:	**$4.6897**	Concurrency:	**177**	Mem BW util:	**23.0%**
Decode tok/s:	**263.8**	Tok/dollar:	**231,887**	Requests:	**708**	Mem util:	**88.0%**
Requests/sec:	**0.515**	$/1M tok:	**4.3124**	Preemptions:	**268**	Nodes:	**3**